

MapReduce Program

Class Text

`org.apache.hadoop.io`

Class Text

[`java.lang.Object`](#)

└ [`org.apache.hadoop.io.BinaryComparable`](#)

└ `org.apache.hadoop.io.Text`

All Implemented Interfaces:

[`Comparable`](#)<[`BinaryComparable`](#)>, [`Writable`](#), [`WritableComparable`](#)<[`BinaryComparable`](#)>

- This class stores text using standard UTF8 encoding.
- It provides methods to serialize, deserialize, and compare texts at byte level.
- They have the “Writable” interface -they know how to write to a DataOutput stream and read from a DataInput stream - explicitly.
- The type of length is integer and is serialized using zero-compressed format.
- In addition, it provides methods for string traversal without converting the byte array to a string.

StubMapper

```
import java.io.IOException;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class StubMapper extends Mapper<LongWritable, Text, Text, IntWritable> {

    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {

        String [] words =value.toString().split(" ");
        for(String word: words)
        {
            context.write(new Text(word), new IntWritable(1));
        }

        /*
        * TODO implement
        */

    }
}
```

StubReducer

```
import java.io.IOException;

import org.apache.hadoop.io.DoubleWritable;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class StubReducer extends Reducer<Text, IntWritable, Text, IntWritable> {

    @Override
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {

        /*
         * TODO implement
         */

        Integer count = 0;
        for(IntWritable val:values)
        {
            count += val.get();
        }
        context.write(key, new IntWritable(count));
    }
}
```

StubDriver

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class StubDriver {

    public static void main(String[] args) throws Exception {

        /*
         * Validate that two arguments were passed from the command line.
         */
        if (args.length != 2) {
            System.out.printf("Usage: StubDriver <input dir> <output dir>\n");
            System.exit(-1);
        }

        /*
         * Instantiate a Job object for your job's configuration.
         */
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "word count");

        /*
         * Specify the jar file that contains your driver, mapper, and reducer.
         * Hadoop will transfer this jar file to nodes in your cluster running
         * mapper and reducer tasks.
         */
        job.setJarByClass(StubDriver.class);
```

StubDriver

```
job.setJarByClass(StubDriver.class);  
  
/*  
 * Specify an easily-decipherable name for the job.  
 * This job name will appear in reports and logs.  
 */  
job.setJobName("Stub Driver");  
  
/*  
 * TODO implement  
 */  
  
job.setMapperClass(StubMapper.class);  
job.setReducerClass(StubReducer.class);  
job.setOutputKeyClass(Text.class);  
job.setOutputValueClass(IntWritable.class);  
  
FileInputFormat.addInputPath(job, new Path(args[0]));  
FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
/*  
 * Start the MapReduce job and wait for it to finish.  
 * If it finishes successfully, return 0. If not, return 1.  
 */  
boolean success = job.waitForCompletion(true);  
System.exit(success ? 0 : 1);
```

StubTester

```
MapDriver<LongWritable, Text, Text, IntWritable> mapDriver;  
ReduceDriver<Text, IntWritable, Text, IntWritable> reduceDriver;  
MapReduceDriver<LongWritable, Text, Text, IntWritable, Text, IntWritable> mapReduceDriver;  
  
/*  
 * Set up the test. This method will be called before every test.  
 */  
@Before  
public void setUp() {  
  
    /*  
     * Set up the mapper test harness.  
     */  
    StubMapper mapper = new StubMapper();  
    mapDriver = new MapDriver<LongWritable, Text, Text, IntWritable>();  
    mapDriver.setMapper(mapper);  
  
    /*  
     * Set up the reducer test harness.  
     */  
    StubReducer reducer = new StubReducer();  
    reduceDriver = new ReduceDriver<Text, IntWritable, Text, IntWritable>();  
    reduceDriver.setReducer(reducer);  
  
    /*  
     * Set up the mapper/reducer test harness.  
     */  
    mapReduceDriver = new MapReduceDriver<LongWritable, Text, Text, IntWritable, Text, IntWritable>();  
    mapReduceDriver.setMapper(mapper);  
}
```

Execution

- Copy input file “pg001.txt” inside the project
- Right click the project and “Run as Application”. Select “Run configurations”. Click on arguments tab and type “pg001.txt output/wordcount” in program arguments.
- Right click the project and “Run as Application”. Select “Java applications”
- After execution is complete, Right click the project and click on refresh to see the new output folder. This will have the results. Delete this folder before executing again.