

Hive Query Language

Hive Query Language (HQL)

1. **Create and manage tables and partitions.**
2. **Support various Relational, Arithmetic, and Logical Operators.**
3. **Evaluate functions.**
4. **Download the contents of a table to a local directory or result of queries to HDFS directory.**

DDL and DML statements

Database

To create a database named “STOCKS” with comments and database properties.

**CREATE DATABASE IF NOT EXISTS STOCKS COMMENT 'STOCK Details' WITH
DBPROPERTIES ('creator' = 'JOHN');**

Database

To display list of databases

SHOW DATABASES;

To describe a database.

DESCRIBE DATABASE STOCKS;

To describe extended database.

DESCRIBE DATABASE EXTENDED STOCKS;

Database

To drop database.

DROP DATABASE STOCKS;

To alter the database properties

ALTER DATABASE STOCKS SET DBPROPERTIES('edited-by'='JAMES')

To make database as the current working database

USE STOCKS

Tables

Hive provides two kinds of tables:

- ▶ **Managed Table**
 - ▶ Hive Stores the table under the warehouse folder under Hive
 - ▶ Complete life cycle of table is managed by Hive
 - ▶ When internal table is dropped, it drops the data as well as the metadata
- ▶ **External Table**
 - ▶ When table is dropped, it retains the data in the underlying location. But drops metadata

Tables

To create managed table named 'STOCK'.

```
CREATE TABLE IF NOT EXISTS STOCK(symbol STRING, reported STRING, price  
FLOAT) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',';
```

To create external table named 'EXT_STOCK'.

```
CREATE EXTERNAL TABLE IF NOT EXISTS EXT_STOCK(symbol STRING, reported  
TIMESTAMP, price FLOAT) ROW FORMAT DELIMITED FIELDS TERMINATED BY ','  
LOCATION '/user/hive/warehouse/stock';
```


Tables

To load data into the table from file named msft-s.csv. Copy msft-s.csv into Downloads folder.

```
LOAD DATA LOCAL INPATH '/home/cloudera/Downloads/msft-s.csv' OVERWRITE  
INTO TABLE EXT_STOCK;
```

To retrieve the student details from “EXT_STOCK” table.

```
SELECT * from EXT_STOCK;
```

Partitions

Partitions

Partitions split the larger dataset into more meaningful chunks.

Hive provides two kinds of partitions: Static Partition and Dynamic Partition.

Use static partition if the value is known at compile time

- To create static partition based on “price” column.

```
CREATE TABLE IF NOT EXISTS STATIC_PART_STOCK (symbol STRING,  
reported TIMESTAMP) PARTITIONED BY (price FLOAT) ROW FORMAT  
DELIMITED FIELDS TERMINATED BY '\t';
```

- Load data into partition table from table.

```
INSERT OVERWRITE TABLE STATIC_PART_STOCK PARTITION (price=27.0)  
SELECT symbol, reported from EXT_STOCK where price=27.0;
```

Partitions

- To create dynamic partition on column date.

```
CREATE TABLE IF NOT EXISTS DYNAMIC_PART_STOCK(symbol STRING, reported  
TIMESTAMP, price FLOAT) PARTITIONED BY (repyear TIMESTAMP) ROW FORMAT  
DELIMITED FIELDS TERMINATED BY '\t';
```

- To load data into a dynamic partition table from table.

```
SET hive.exec.dynamic.partition = true;
```

Note: The dynamic partition strict mode requires at least one static partition column. To turn this off,
set `hive.exec.dynamic.partition.mode=nonstrict`

```
INSERT OVERWRITE TABLE DYNAMIC_PART_STOCK PARTITION (repyear)  
SELECT symbol,reported, price, YEAR(reported) AS repyear from EXT_STOCK;
```

Bucketing In Hive

- ▶ Partitioning based on geographic locations like country, some bigger countries will have large partitions (ex: 4-5 countries itself contributing 70-80% of total data) where as small countries data will create small partitions (remaining all countries in the world may contribute to just 20-30 % of total data).
- ▶ To overcome the problem of over partitioning, Hive provides Bucketing concept.
- ▶ Features
 - ▶ Bucketing concept is based on (**hashing function on the bucketed column**) **mod (by total number of buckets)**. The hash_function depends on the type of the bucketing column.
 - ▶ Records with the same bucketed column will always be stored in the same bucket.
 - ▶ We use **CLUSTERED BY** clause to divide the table into buckets.
 - ▶ Physically, each bucket is just a file in the table directory, and Bucket numbering is 1-based.
 - ▶ Bucketed tables will create almost equally distributed data file parts.

Advantages of Bucketing

- ▶ Bucketed tables offer efficient sampling. With sampling, we can try out queries on a fraction of data for testing and debugging purpose when the original data sets are very huge.
- ▶ As the data files are equal sized parts, map-side joins will be faster on bucketed tables than non-bucketed tables. In Map-side join, a mapper processing a bucket of the left table knows that the matching rows in the right table will be in its corresponding bucket, so it only retrieves that bucket (which is a small fraction of all the data stored in the right table).
- ▶ Similar to partitioning, bucketed tables provide faster query responses than non-bucketed tables.
- ▶ Bucketing concept also provides the flexibility to keep the records in each bucket to be sorted by one or more columns. This makes map-side joins even more efficient, since the join of each bucket becomes an efficient merge-sort.

Buckets

set hive.enforce.bucketing=true;

- To create a bucketed table having 3 buckets.

**CREATE TABLE IF NOT EXISTS STOCK_BUCKET(symbol STRING,
reported TIMESTAMP, price FLOAT)
CLUSTERED BY (price) into 3 buckets;**

- Load data to bucketed table.

**FROM EXT_STOCK
INSERT OVERWRITE TABLE STOCK_BUCKET
SELECT symbol, reported, price;**

- To display the content of first bucket.

**SELECT DISTINCT price FROM STOCK_BUCKET
TABLESAMPLE(BUCKET 1 OUT OF 3 ON price);**

Aggregations

Hive supports aggregation functions like avg, count, etc.

To write the average and count aggregation function.

```
SELECT avg(price) FROM EXT_STOCK;
```

```
SELECT count(*) FROM EXT_STOCK;
```


Group by and Having

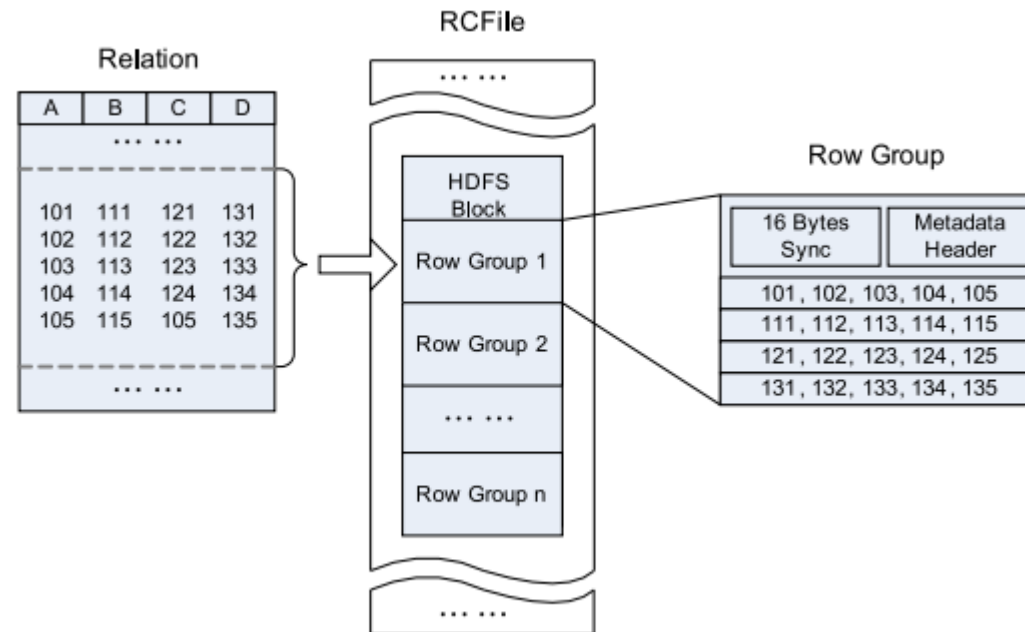
To write group by and having function.

```
SELECT symbol, reported, price  
FROM EXT_STOCK  
GROUP BY symbol, reported, price  
HAVING price > 27.0;
```

Hive File Format

- **Text File** - The default file format is text file.
- **Sequential File** - Sequential files are flat files that store binary key-value pairs.
- **RCFile (Record Columnar File)** - RCFile stores the data in **Column Oriented Manner** which ensures that **Aggregation** operation is not an expensive operation. RCFile first partitions horizontally and then vertically to serialize data.

```
CREATE TABLE STOCK_RC (symbol STRING, reported TIMESTAMP, price FLOAT) STORED AS RCFILE;  
INSERT OVERWRITE TABLE STOCK_RC SELECT symbol, reported, price FROM STOCK;  
SELECT SUM(price) FROM STOCK_RC;
```



Complex Types

```
CREATE TABLE complex (
```

```
  c1 ARRAY<INT>,
```

```
  c2 MAP<STRING, INT>,
```

```
  c3 STRUCT<a:STRING, b:INT, c:DOUBLE>);
```

```
SELECT c1[0], c2['b'], c3.c FROM complex;
```

SerDer

- SerDer stands for Serializer/Deserializer.
- Contains the logic to convert unstructured data into records.
- Implemented using Java.
- Serializers are used at the time of writing. Serializer takes java object and translates it into something Hive can write to HDFS
- Deserializers are used at query time (SELECT Statement). Deserializer converts binary representation or a string of record to java object that Hive can manipulate

```
CREATE TABLE stations (usaf STRING, wban STRING, name STRING) ROW FORMAT  
SERDE 'org.apache.hadoop.hive.contrib.serde2.RegexSerDe' WITH  
SERDEPROPERTIES ("input.regex" = "(\\d{6}) (\\d{5}) (.{29}) .*");
```

Manipulate XML Data

```
CREATE TABLE XMLSAMPLE(xmldata string);
```

```
LOAD DATA LOCAL INPATH '/home/cloudera/Downloads/employee.xml' INTO TABLE  
XMLSAMPLE
```

```
CREATE TABLE xpath_table AS  
SELECT xpath_int(xmldata, 'employee/empid'),  
Xpath_string(xmldata,'employee/name')  
FROM xmlsample;
```

```
SELECT * FROM xpath_table;
```

Further Readings

- ▶ <http://en.wikipedia.org/wiki/RCFile>
- ▶ <https://cwiki.apache.org/confluence/display/Hive/DynamicPartitions>
- ▶ <https://cwiki.apache.org/confluence/display/Hive/LanguageManual+DDL>
- ▶ <https://cwiki.apache.org/confluence/display/Hive/LanguageManual+DML>